



효소 열안정성 예측

Enzyme Stability Prediction

7조 김효진, 나다경, 안이현, 유도현



Contents

01 About competition

02 Feature Engineering

03 Modeling

04 Results



About competition

Goal, Data, EDA

Goal of competition

Test set

| seq_id | protein_sequence | tm (melting temperature) | rank |
|--------|------------------|-----------------------------|------|
| 1 | VPVNEPD ... | 53.6 | 483 |
| 2 | VPVNPAPD ... | 61.2 | 318 |

Our prediction

Wild type enzyme :



Enzyme variants :
One point mutation



[deletion]



[substitution]

Data

Train set

| | protein_sequence | pH | data_source | tm |
|--------|---|-----|-----------------------------------|------|
| seq_id | | | | |
| 0 | AAAAKAAALALLGEAPEVVDIWLPAGWRQPFRVFRLERKGDGVLVG... | 7.0 | doi.org/10.1038/s41592-020-0801-4 | 75.7 |

Melting temperature

28981 rows × 5 columns

Test set

```
test_enzyme = "VPVNPEPDATSVENVALKTGSGDSQSDPIKADLEVKGQSALPFDVDCWAILCKGAPNVLQRVNEKTKNSNRDRSGANKGPFKDPQKWGIKALPPKNPSWSAQDFKSPEE  
YAFASSLQGGTNAILAPVNLASQNSQGGVLNGFYANKVAQFDPSKPQQTGTWTFQITKFTGAAGPYCKALGSNDKSVCDKNKNIAGDWGFDPKWAYQYDEKNNKFNYVGK"
```

| | protein_sequence | pH | data_source |
|--------|---|----|-------------|
| seq_id | | | |
| 31390 | VPVNPEPDATSVENVAEKTGSGDSQSDPIKADLEVKGQSALPFDVD... | 8 | Novozymes |
| 31391 | VPVNPEPDATSVENVAKKTGSGDSQSDPIKADLEVKGQSALPFDVD... | 8 | Novozymes |
| 31392 | VPVNPEPDATSVENVAKTGSGDSQSDPIKADLEVKGQSALPFDVDC... | 8 | Novozymes |

2413 rows × 4 columns

One point mutation


→ We can check where the mutation occurred.

PDB file

```
!pip install py3Dmol -q
import py3Dmol
with open("../input/novozymes-enzyme-stability-prediction/wildtype_structure_prediction_af2.pdb") as ifile:
    protein = "".join([x for x in ifile])
#view = py3Dmol.view(query='pdb:1DIV', width=800, height=600)
view = py3Dmol.view(width=800, height=600)
view.addModelsAsFrames(protein)
style = {'cartoon': {'color': 'spectrum'}, 'stick': {}}
view.setStyle({'model': -1}, style)
view.zoom(0.12)
view.rotate(235, {'x':0, 'y':1, 'z':1})
view.spin({'x':-0.2, 'y':0.5, 'z':1}, 1)
view.show()
```



➡ Provide protein structure information



B-factor & pLDDT

Alphafold

- The best AI to predict protein structure.
- Alphafold provides a thermal stability feature called plddt.
- In general, it is known that the higher the plddt, the higher the thermal stability.

B-factor

Tm

pLDDT

B-factor

- one of the protein properties provided in the "original" pdb file. (Not in our file)
- B-factor is an indicator of thermal motion about atom.
- It has a high correlation with our target, tm.

We can use pLDDT instead of B-factor.

EDA

| | Train | Test |
|------------------|----------------------|----------------------|
| Protein sequence | Wild type + mutation | One – point mutation |
| pH | 1-14 | 8 (fixed) |
| Data source | diverse | fixed |
| Sequence length | diverse | fixed(220 or 221) |



Create new Train set similar to the Test set!



Feature Engineering

Adding features

Create new train data

Original Train set

| | protein_sequence | pH | data_source | tm |
|--------|--|-----|-----------------------------------|------|
| seq_id | | | | |
| 0 | AAAAKAAALALLGEAPEVVDIWLPGWRQPFRVFRLERKGDGVLVG... | 7.0 | doi.org/10.1038/s41592-020-0801-4 | 75.7 |

Similar
enzyme
groups

Wildtype
(original)
Amino acid

Mutated
Amino
acid

dTm
-> New target!



Using PDB file + fixing pH level


| | PDB | WT | position | MUT | dTm | sequence | mutant_seq |
|---|------|----|----------|-----|---------|---|--|
| 0 | GP01 | L | 89 | A | 2.28642 | MLVMTEYLLSAGICMAIVSILLIGMAISNVSKGQYAKRFFFFATSC... | MLVMTEYLLSAGICMAIVSILLIGMAISNVSKGQYAKR |
| 1 | GP01 | T | 95 | C | 1.48642 | MLVMTEYLLSAGICMAIVSILLIGMAISNVSKGQYAKRFFFFATSC... | MLVMTEYLLSAGICMAIVSILLIGMAISNVSKGQYAKR |
| 2 | GP01 | T | 95 | C | 0.28642 | MLVMTEYLLSAGICMAIVSILLIGMAISNVSKGQYAKRFFFFATSC... | MLVMTEYLLSAGICMAIVSILLIGMAISNVSKGQYAKR |
| 3 | GP01 | T | 95 | S | 2.48642 | MLVMTEYLLSAGICMAIVSILLIGMAISNVSKGQYAKRFFFFATSC... | MLVMTEYLLSAGICMAIVSILLIGMAISNVSKGQYAKR |
| 4 | GP01 | T | 95 | S | 3.88642 | MLVMTEYLLSAGICMAIVSILLIGMAISNVSKGQYAKRFFFFATSC... | MLVMTEYLLSAGICMAIVSILLIGMAISNVSKGQYAKR |

Adding domain knowledge data

```
aa_props = pd.read_csv('../input/aminoacids-physical-and-chemical-properties/amin  
oacids.csv').set_index('Letter')  
PROPS = ['Molecular Weight', 'Residue Weight', 'pKa1', 'pKb2', 'pKx3', 'pI4',  
         'H', 'VSC', 'P1', 'P2', 'SASA', 'NCISC']  
print('Amino Acid properties dataframe. Shape:', aa_props.shape )  
aa_props.head(22)
```

| | Name | Abbr | Molecular Weight | Molecular Formula | Residue Formula | Residue Weight | pKa1 | pKb2 | pKx3 | pI4 |
|--------|----------------|------|------------------|-------------------|-----------------|----------------|------|-------|-------|------|
| Letter | | | | | | | | | | |
| A | Alanine | Ala | 89.10 | C3H7NO2 | C3H5NO | 71.08 | 2.34 | 9.69 | NaN | 6.00 |
| C | Cysteine | Cys | 121.16 | C3H7NO2S | C3H5NOS | 103.15 | 1.96 | 10.28 | 8.18 | 5.07 |
| D | Aspartic acid | Asp | 133.11 | C4H7NO4 | C4H5NO3 | 115.09 | 1.88 | 9.60 | 3.65 | 2.77 |
| E | Glutamic acid | Glu | 147.13 | C5H9NO4 | C5H7NO3 | 129.12 | 2.19 | 9.67 | 4.25 | 3.22 |
| F | Phenylalanine | Phe | 165.19 | C9H11NO2 | C9H9NO | 147.18 | 1.83 | 9.13 | NaN | 5.48 |
| G | Glycine | Gly | 75.07 | C2H5NO2 | C2H3NO | 57.05 | 2.34 | 9.60 | NaN | 5.97 |
| H | Histidine | His | 155.16 | C6H9N3O2 | C6H7N3O | 137.14 | 1.82 | 9.17 | 6.00 | 7.59 |
| I | Isoleucine | Ile | 131.18 | C6H13NO2 | C6H11NO | 113.16 | 2.36 | 9.60 | NaN | 6.02 |
| K | Lysine | Lys | 146.19 | C6H14N2O2 | C6H12N2O | 128.18 | 2.18 | 8.95 | 10.53 | 9.74 |
| L | Leucine | Leu | 131.18 | C6H13NO2 | C6H11NO | 113.16 | 2.36 | 9.60 | NaN | 5.98 |
| M | Methionine | Met | 149.21 | C5H11NO2S | C5H9NOS | 131.20 | 2.28 | 9.21 | NaN | 5.74 |
| N | Asparagine | Asn | 132.12 | C4H8N2O3 | C4H6N2O2 | 114.11 | 2.02 | 8.80 | NaN | 5.41 |
| O | Hydroxyproline | Hyp | 131.13 | C5H9NO3 | C5H7NO2 | 113.11 | 1.82 | 9.65 | NaN | NaN |
| P | Proline | Pro | 115.13 | C5H9NO2 | C5H7NO | 97.12 | 1.99 | 10.60 | NaN | 6.30 |
| Q | Glutamine | Gln | 146.15 | C5H11NO2 | C5H9NO | 129.12 | 2.19 | 9.67 | 4.25 | 3.22 |

Transformer ESM features + embeddings



PCA

```
ESM2(  
  (embed_tokens): Embedding(33, 1280, padding_idx=1)
```

1280 : large dimension

```
  (layers): ModuleList(  
    (0): TransformerLayer(  
      (self_attn): MultiheadAttention(  
        (k_proj): Linear(in_features=1280, out_features=1280, bias=True)  
        (v_proj): Linear(in_features=1280, out_features=1280, bias=True)  
        (q_proj): Linear(in_features=1280, out_features=1280, bias=True)  
        (out_proj): Linear(in_features=1280, out_features=1280, bias=True)  
        (rot_emb): RotaryEmbedding()      )  
      (self_attn_layer_norm): LayerNorm1D()  
      (fc1): Linear(in_features=1280, out_features=1280, bias=True)  
      (fc2): Linear(in_features=1280, out_features=1280, bias=True)  
      (final_layer_norm): LayerNorm1D()  
    )  
  )  
)
```

reduce dimension
1280 to 32

REDUCE EMBEDDING DIM FROM 1280 TO 32 OR 16 WITH PCA

```
from cuml import PCA
```

```
pca_pool = PCA(n_components=32)
```

```
pca_embeds = pca_pool.fit_transform(all_pdb_embed_pool.astype('float32'))
```

```
pca_local = PCA(n_components=16)
```

```
pca_local.fit(all_pdb_embed_tmp.astype('float32'))
```

```
pdb_map = {x:y for x,y in zip(all_pdb, range(len(all_pdb)))}
```

```
pdb_map['kaggle'] = len(all_pdb)
```


```
del all_pdb_embed_tmp
```

```
_ = gc.collect()
```


Feature engineering

VPVNPEPDATSVENVALKTGSGDSQSDPIKADLEVKGQSALPFDVDCWAILCKGAPN
VLQRVNEKTKNSNRDRSGANKGPFKDPQKWGIKALPPKNPSWSAQDFKSPEEYAFAS
SLQGGTNAILAPVNLASQNSQGGVNLNGFYSA NKVAQFDP5KPOQT KGTW FQJTKFTG
AAGPYCKALGSNDKSVCDKNKNIAGDWGFDPKWAYQYDEKNNKFNYVGK

Wild Type & Mutant Sequences



Wild Type Structure information(pLDDT)



Amino acids properties



Transformer ESM + Embeddings


Extracting embeddings from proteins...

GP01 , GP02 , GP03 , GP04 , GP06 , GP07 , GP08 , GP09 , GP10 , GP11 , GP12 , GP13
 , GP14 , GP15 , GP16 , GP17 , GP18 , GP19 , GP20 , GP21 , GP22 , GP23 , GP24 , GP
25 , GP26 , GP27 , GP28 , GP29 , GP30 , GP31 , GP32 , GP33 , GP34 , GP35 , GP36 ,
GP37 , GP38 , GP39 , GP40 , GP41 , GP42 , GP43 , GP44 , GP45 , GP46 , GP48 , GP49




PCA Model

127 features



Modeling



Predict dTM



Modeling

XGBoost, LGBM, Random Forest, Ridge

Models

Data set

127 features

⇒ The problem of overfitting

⇒ We have to solve this problem!



Models

XGBoost

LGBM

Randomforest


Ridge

Feature importance

LGBM Feature Importance - Top10



XGB GAIN Feature Importance - Top 10




<Best features>

1. P1_delta
2. Pca_mutant
3. Pca_pool
4. Pca_local
5. pca_wt


⇒ the polarity difference of a single amino acid has a high effect on thermal stability.

⇒ the overall structure of the protein has a significant impact on thermal stability

K-fold vs Group K-fold



Randomly divide the fold




- Each group appears **Once** across all folds
- ⇒ The same group is **Not** represented in both testing / validation and training sets
- ⇒ make it possible to **detect overfitting** situations

Comparison

XGBoost

```
model = xgb.train(xgb_parms,  
                  dtrain=dtrain,  
                  evals=[(dtrain, 'train'), (dvalid, 'valid')],  
                  num_boost_round=9999,  
                  early_stopping_rounds=100,  
                  verbose_eval=100)  
model.save_model(f'xgb_models/XGB_fold{fold}.xgb')
```

Test preds histogram




Spearman Metric : 0.341,
Leader Board Score : 0.273

LGBM

```
model = lgb.train(params,  
                  dtrain,  
                  valid_sets=[dtrain, dvalid],  
                  early_stopping_rounds=10)  
#model.save_model(f'lgbm_models/LGBM_fold{fold}.lgb')  
joblib.dump(model, f'lgb_{fold}.pkl')
```

Test preds histogram



Spearman Metric : 0.346,
Leader Board Score : 0.291

Comparison


Ridge

```
params_ridge={'alpha':[0.001, 0.01, 0.1, 1, 10, 100, 1000]}
gscv_ridge = GridSearchCV(model, param_grid=params_ridge,
                           scoring="neg_root_mean_squared_error",
                           n_jobs=-1,
                           cv=skf)

gscv_ridge.fit(X, y, groups=train.group)

gscv_ridge.best_params_
model.fit(X_train, y_train)
```

Test preds histogram




Spearman Metric : 0.331
Leader Board Score : 0.255

Randomforest

```
model = RandomForestRegressor(n_estimators=200,
                              bootstrap=True,
                              max_depth=8,
                              min_samples_split=4,
                              min_samples_leaf=5,
                              max_features=12,
                              random_state=SEED)

model.fit(X_train, y_train)
```

Test preds histogram



Spearman Metric : 0.341,
Leader Board Score : 0.273



Results

Submission

Best model

| Model | XGBoost | LGBM | RandomForest | Ridge |
|-------------------------------------|---------|--------|--------------|-------|
| RMSE | 0.269 | 0.2704 | 0.272 | 0.278 |
| Spearman Metric (predicting dTm) | 0.341 | 0.3455 | 0.348 | 0.331 |
| Leader Board Score | 0.273 | 0.291 | 0,264 | 0.255 |

| | seq_id | tm | rank |
|---|--------|----------|--------|
| 0 | 31390 | 0.504586 | 2167.0 |
| 1 | 31391 | 0.511375 | 1512.0 |
| 2 | 31392 | 0.511383 | 1472.0 |
| 3 | 31393 | 0.517814 | 95.0 |
| 4 | 31394 | 0.515988 | 414.0 |



Thank You